

Image Representation, Processing, and Computer Vision

Ajad Chhatkuli and Bishesh Khanal

Dec 23, 2018

Contents

Human vs. Computer Vision





- ▶ What do we see ?



- ▶ What do we see ?
- ▶ Where are those located ?



- ▶ What do we see ?
- ▶ Where are those located ?
- ▶ ...

Image Formation



What we see

Image Formation



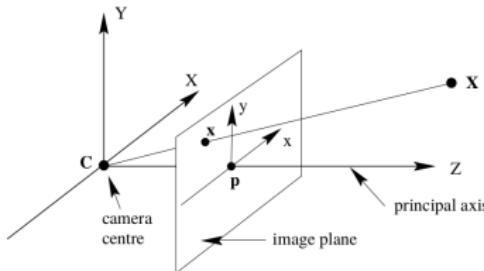
What we see

What machines see

Image Formation



What we see



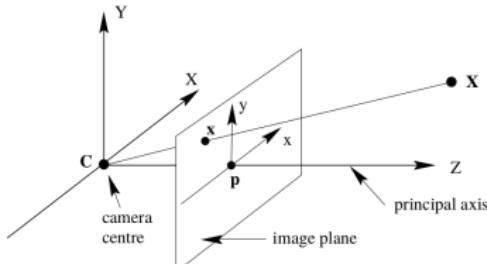
Source: Hartley, Richard, and Andrew Zisserman. *Multiple view geometry in computer vision*. Cambridge university press, 2003.

What machines see

Image Formation



What we see



Source: Hartley, Richard, and Andrew Zisserman. Multiple view geometry in computer vision. Cambridge university press, 2003.

What machines see

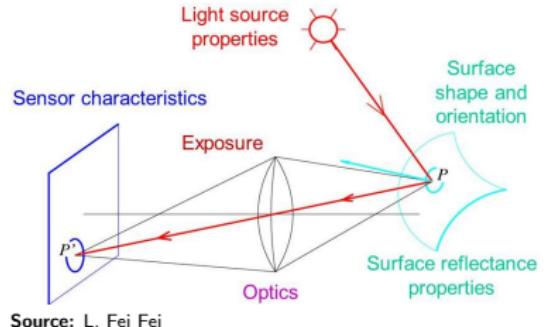


Image: Function of Multiple Spatial Variables

- ▶ Computers just see numbers (in fact, a lot of numbers)
- ▶ 2D image: each discrete position (x, y) consist of intensity values
- ▶ 3D image: each discrete position (x, y, z) consist of intensity values
- ▶ grayscale images represented as $N \times M$ matrix
- ▶ color images (**R**, **G**, **B**) represented as $3 \times N \times M$ matrix

Two Closest images ?



Distance Between Images ?

Distance Between Images ?

$$A = \begin{pmatrix} 2 & 3 & 4 \\ 1 & 7 & 9 \end{pmatrix}$$

$$\text{Vec}(A) = \begin{pmatrix} 2 \\ 3 \\ 4 \\ 1 \\ 7 \\ 9 \end{pmatrix}$$

Distance Between Images ?

$$A = \begin{pmatrix} 2 & 3 & 4 \\ 1 & 7 & 9 \end{pmatrix}$$

$$\text{Vec}(A) = \begin{pmatrix} 2 \\ 3 \\ 4 \\ 1 \\ 7 \\ 9 \end{pmatrix}$$

Remember distance between vectors in Euclidean space ?

Distance Between Images ?

$$A = \begin{pmatrix} 2 & 3 & 4 \\ 1 & 7 & 9 \end{pmatrix}$$

$$\text{Vec}(A) = \begin{pmatrix} 2 \\ 3 \\ 4 \\ 1 \\ 7 \\ 9 \end{pmatrix}$$

Remember distance between vectors in Euclidean space ?

Sum of squared distance:

$$\sum ((x_1 - y_1)^2 + (x_2 - y_2)^2 + \dots + (x_n - y_n)^2)$$

Distance Between Images ?

$$A = \begin{pmatrix} 2 & 3 & 4 \\ 1 & 7 & 9 \end{pmatrix}$$

$$\text{Vec}(A) = \begin{pmatrix} 2 \\ 3 \\ 4 \\ 1 \\ 7 \\ 9 \end{pmatrix}$$

Remember distance between vectors in Euclidean space ?

Sum of squared distance:

$$\sum ((x_1 - y_1)^2 + (x_2 - y_2)^2 + \dots + (x_n - y_n)^2)$$

Can you think of other distances ?

Distance Between Images ?

$$A = \begin{pmatrix} 2 & 3 & 4 \\ 1 & 7 & 9 \end{pmatrix}$$

$$\text{Vec}(A) = \begin{pmatrix} 2 \\ 3 \\ 4 \\ 1 \\ 7 \\ 9 \end{pmatrix}$$

Remember distance between vectors in Euclidean space ?

Sum of squared distance:

$$\sum ((x_1 - y_1)^2 + (x_2 - y_2)^2 + \dots + (x_n - y_n)^2)$$

Can you think of other distances ?

Sum of absolute values of differences:

$$\sum (|x_1 - y_1| + |x_2 - y_2| + \dots + |x_n - y_n|)$$

Distance Between Images ?

$$A = \begin{pmatrix} 2 & 3 & 4 \\ 1 & 7 & 9 \end{pmatrix}$$

$$\text{Vec}(A) = \begin{pmatrix} 2 \\ 3 \\ 4 \\ 1 \\ 7 \\ 9 \end{pmatrix}$$

Remember distance between vectors in Euclidean space ?

Sum of squared distance:

$$\sum ((x_1 - y_1)^2 + (x_2 - y_2)^2 + \dots + (x_n - y_n)^2)$$

Can you think of other distances ?

Sum of absolute values of differences:

$$\sum (|x_1 - y_1| + |x_2 - y_2| + \dots + |x_n - y_n|)$$

Norms: L_1, L_2, \dots

Two Closest images ?



Two Closest images ?



Based on SSD, Kitten images are further apart!

Two Closest images ?



Based on SSD, Kitten images are further apart!
Euclidean space not suitable in most cases for images

Two Closest images ?



Based on SSD, Kitten images are further apart!
Euclidean space not suitable in most cases for images
Better representation ? what manifold do the images lie ?
Better metric “distance”

Dimensionality Reduced Representation, Feature Extraction



Dimensionality Reduced Representation, Feature Extraction



- ▶ An image of 600×400 size is a point in 240000 dimensional space

Dimensionality Reduced Representation, Feature Extraction



- ▶ An image of 600×400 size is a point in 240000 dimensional space
- ▶ Even with all the computational might that we have now, this high dimensional data is **too large** to process

Dimensionality Reduced Representation, Feature Extraction



- ▶ An image of 600×400 size is a point in 240000 dimensional space
- ▶ Even with all the computational might that we have now, this high dimensional data is **too large** to process
- ▶ Can extract a small set of **informative** and **non-redundant** features

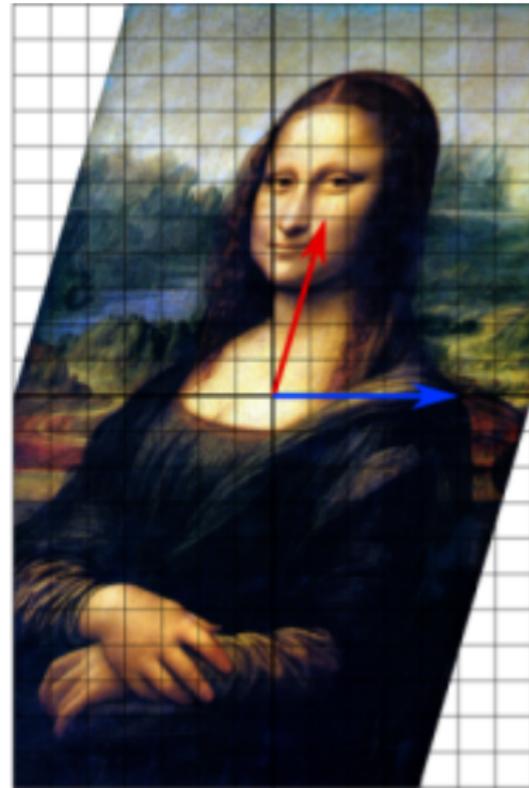
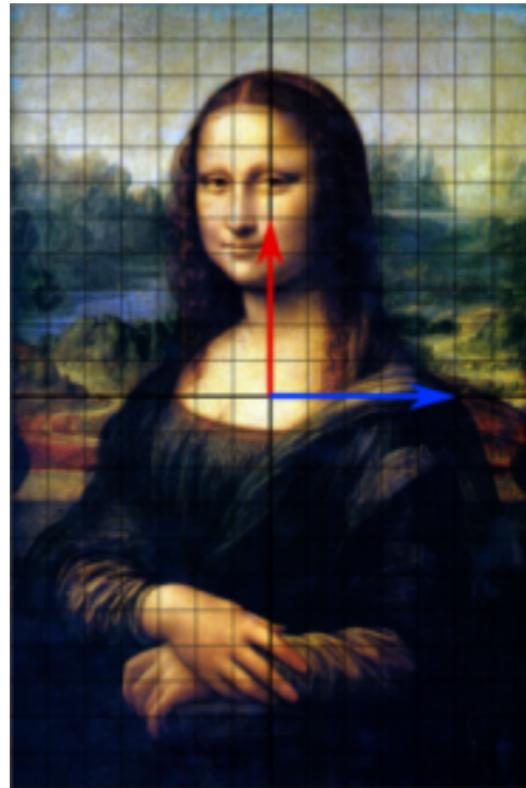
Dimensionality Reduced Representation, Feature Extraction



- ▶ An image of 600×400 size is a point in 240000 dimensional space
- ▶ Even with all the computational might that we have now, this high dimensional data is **too large** to process
- ▶ Can extract a small set of **informative** and **non-redundant** features
- ▶ Can represent image differently in a low dimensional space

Principal Component Analysis

Eigenvectors



Source: By TreyGreer62 <https://commons.wikimedia.org/w/index.php?curid=12768508>

Eigenvectors and Eigenvalues

A matrix A acts on vector \vec{x} to give a new vector \vec{b}

$$\text{E.g. } A = \begin{bmatrix} 3 & -1 \\ -1 & 3 \end{bmatrix} \quad \vec{x} = \begin{bmatrix} 1 \\ 2 \end{bmatrix} \quad A\vec{x} = \vec{b} = \begin{bmatrix} 1 \\ 5 \end{bmatrix}$$

- ▶ Generally, \vec{x} and \vec{b} are in different direction

Eigenvectors and Eigenvalues

A matrix A acts on vector \vec{x} to give a new vector \vec{b}

$$\text{E.g. } A = \begin{bmatrix} 3 & -1 \\ -1 & 3 \end{bmatrix} \quad \vec{x} = \begin{bmatrix} 1 \\ 2 \end{bmatrix} \quad A\vec{x} = \vec{b} = \begin{bmatrix} 1 \\ 5 \end{bmatrix}$$

- ▶ Generally, \vec{x} and \vec{b} are in different direction
- ▶ If $\vec{x} = \begin{bmatrix} 1 \\ 1 \end{bmatrix}$, $\vec{b} = \begin{bmatrix} 2 \\ 2 \end{bmatrix} = 2\vec{x}$
- ▶ $\begin{bmatrix} 1 \\ 1 \end{bmatrix}$ and $\begin{bmatrix} 2 \\ 2 \end{bmatrix}$ in same direction

Eigenvectors and Eigenvalues

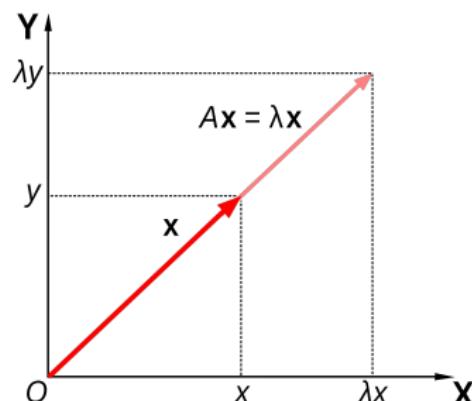
A matrix A acts on vector \vec{x} to give a new vector \vec{b}

$$\text{E.g. } A = \begin{bmatrix} 3 & -1 \\ -1 & 3 \end{bmatrix} \quad \vec{x} = \begin{bmatrix} 1 \\ 2 \end{bmatrix} \quad A\vec{x} = \vec{b} = \begin{bmatrix} 1 \\ 5 \end{bmatrix}$$

- ▶ Generally, \vec{x} and \vec{b} are in different direction
- ▶ If $\vec{x} = \begin{bmatrix} 1 \\ 1 \end{bmatrix}$, $\vec{b} = \begin{bmatrix} 2 \\ 2 \end{bmatrix} = 2\vec{x}$
- ▶ $\begin{bmatrix} 1 \\ 1 \end{bmatrix}$ and $\begin{bmatrix} 2 \\ 2 \end{bmatrix}$ in same direction

If $A\vec{x} = \lambda\vec{x}$,

- ▶ \vec{x} : eigenvector and λ : eigenvalue



Source: Lyudmil Antonov Lantonov
https://en.wikipedia.org/wiki/Eigenvalues_and_eigenvectors

Principal Component Analysis

Let's work through a code!

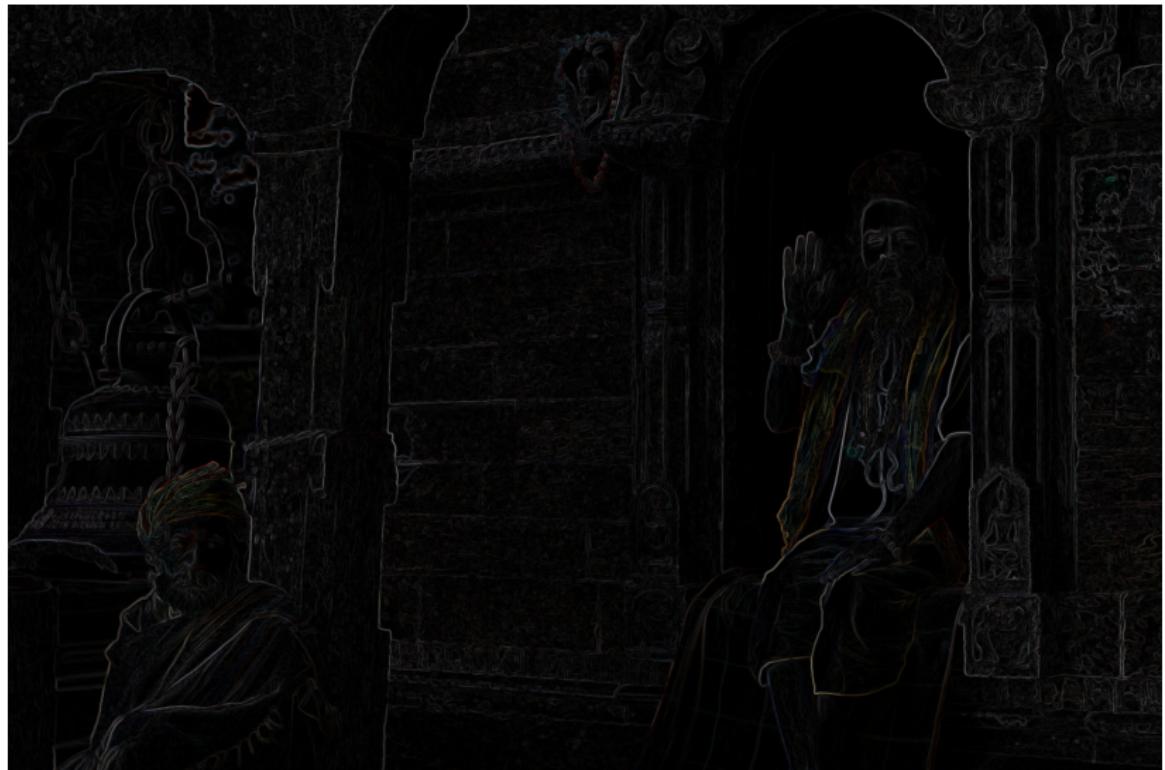
Feature Extraction : Hand Engineered

- ▶ Edge detection
- ▶ Shape detection
- ▶ Thresholding
- ▶ SIFT
- ▶ ...

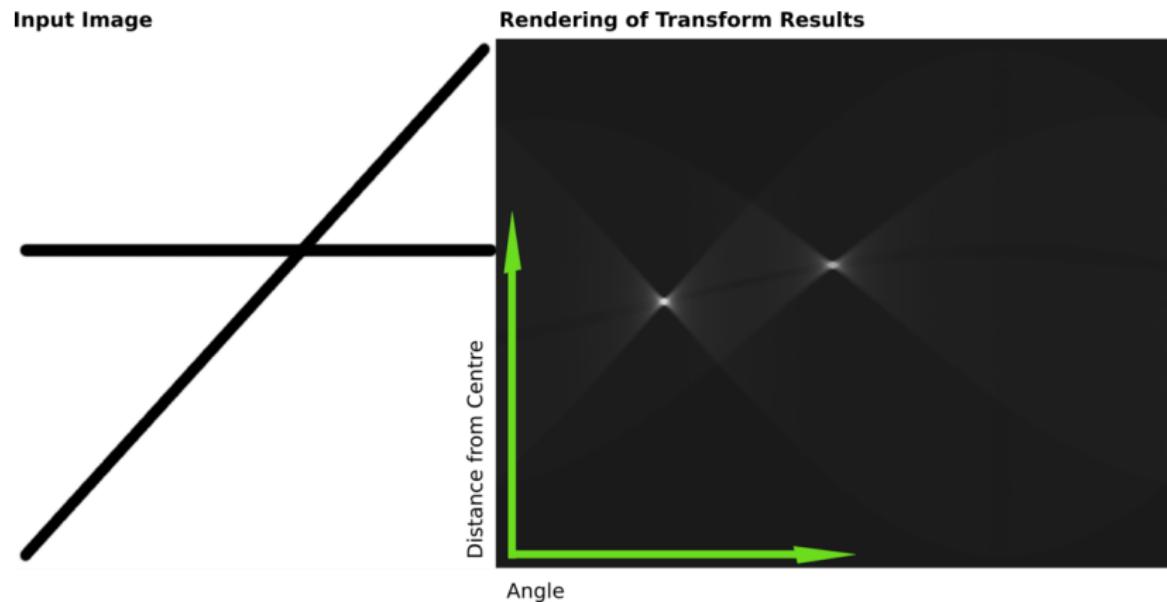
Feature Extraction : Hand Engineered - Edge detection



Feature Extraction : Hand Engineered - Edge detection



Feature Extraction : Hand Engineered - Shape detection



source: https://en.wikipedia.org/wiki/Hough_transform

Feature Extraction : Hand Engineered - Thresholding



Feature Extraction : Hand Engineered - SIFT



source: https://en.wikipedia.org/wiki/Scale-invariant_feature_transform

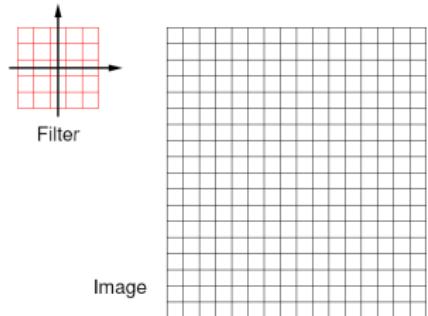
Linear Filters, Correlation and Convolution

Let $I(X, Y)$ a $n \times n$ image

Let $F(X, Y)$ a $m \times m$ image ("filter")

- Here $m = 5$. Let $k = m/2 = 2$

Compute a new image as:



$$I'(X, Y) = \sum_{j=-k}^k \sum_{i=-k}^k F(i, j) I(X + i, Y + j)$$

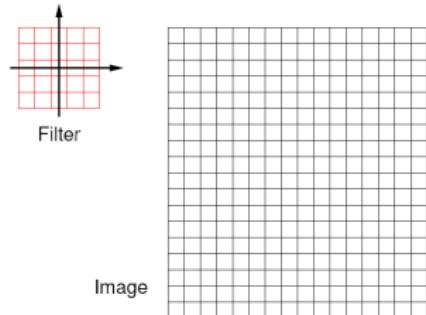
Linear Filters, Correlation and Convolution

Let $I(X, Y)$ a $n \times n$ image

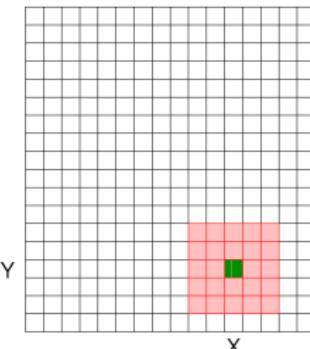
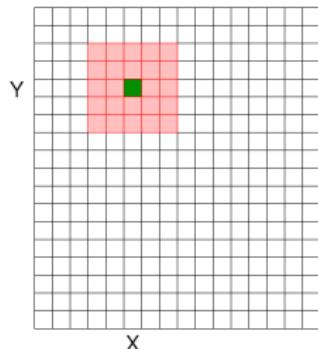
Let $F(X, Y)$ a $m \times m$ image ("filter")

- Here $m = 5$. Let $k = m/2 = 2$

Compute a new image as:



$$I'(X, Y) = \sum_{j=-k}^k \sum_{i=-k}^k F(i, j) I(X + i, Y + j)$$



Source: Bob Woodham, UBC. CPSC
425: Computer Vision lecture notes

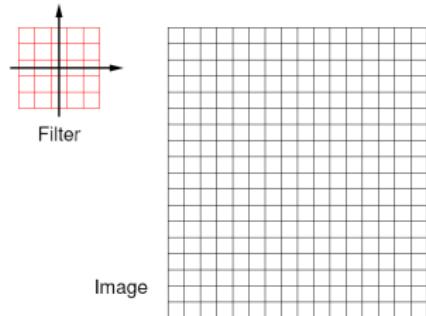
Linear Filters, Correlation and Convolution

Let $I(X, Y)$ a $n \times n$ image

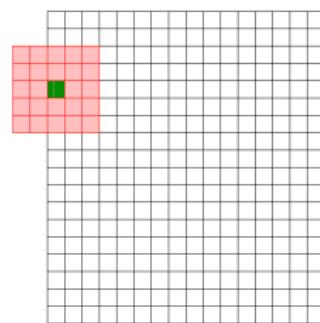
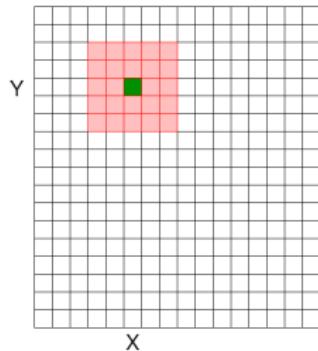
Let $F(X, Y)$ a $m \times m$ image ("filter")

- Here $m = 5$. Let $k = m/2 = 2$

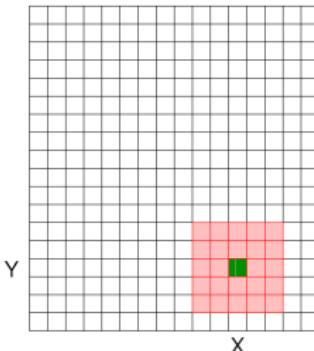
Compute a new image as:



$$I'(X, Y) = \sum_{j=-k}^k \sum_{i=-k}^k F(i, j) I(X + i, Y + j)$$



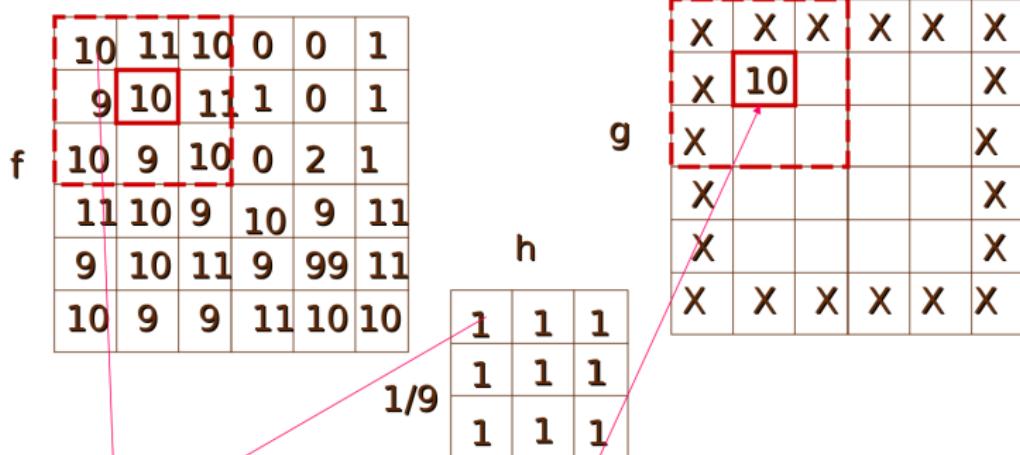
Boundary effects



Source: Bob Woodham, UBC. CPSC
425: Computer Vision lecture notes

Computation Example

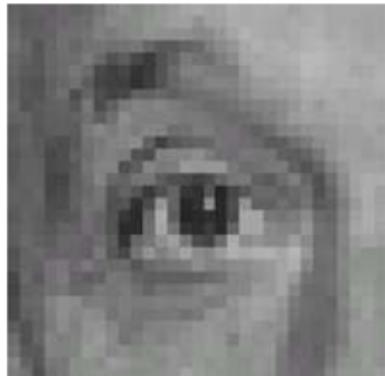
$$g(X, Y) = \sum_{j=-1}^1 \sum_{i=-1}^1 h(i, j)f(X + i, Y + j)$$



$$\begin{aligned} & 1/9.(10 \times 1 + 11 \times 1 + 10 \times 1 + 9 \times 1 + 10 \times 1 + 11 \times 1 + 10 \times 1 \\ & + 9 \times 1 + 10 \times 1) = 1/9.(90) = 10 \end{aligned}$$

Source: Jana Kosecka, Stanford CS223b Computer Vision, Winter 2007

Linear Filtering



Original

| | | |
|---|---|---|
| 0 | 0 | 0 |
| 0 | 1 | 0 |
| 0 | 0 | 0 |

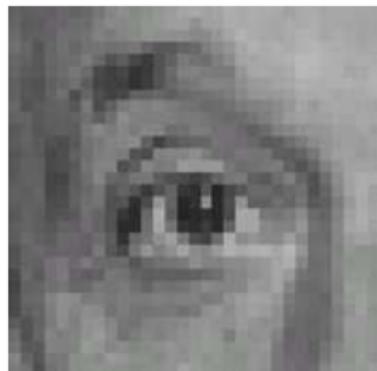
Filter

?

Result

Source: Bob Woodham, UBC. CPSC 425: Computer Vision lecture notes

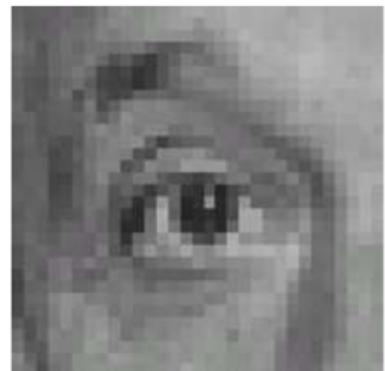
Linear Filtering



Original

| | | |
|---|---|---|
| 0 | 0 | 0 |
| 0 | 1 | 0 |
| 0 | 0 | 0 |

Filter

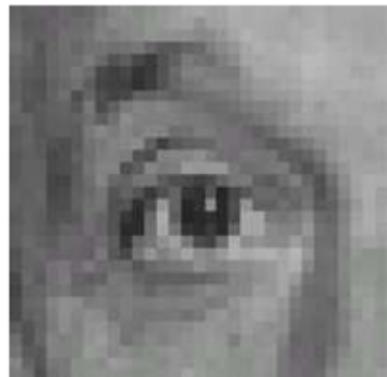


Result

(no change)

Source: Bob Woodham, UBC. CPSC 425: Computer Vision lecture notes

Linear Filtering



$$\frac{1}{9}$$

| | | |
|---|---|---|
| 1 | 1 | 1 |
| 1 | 1 | 1 |
| 1 | 1 | 1 |

?

Original

Filter

Result

(filter sums to 1)

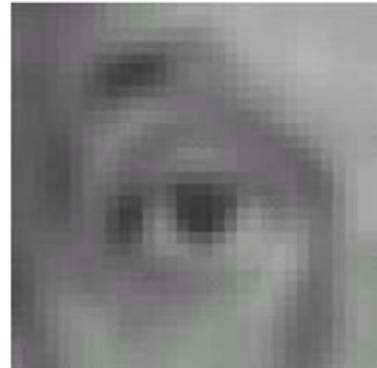
Source: Bob Woodham, UBC. CPSC 425: Computer Vision lecture notes

Linear Filtering



$$\frac{1}{9}$$

| | | |
|---|---|---|
| 1 | 1 | 1 |
| 1 | 1 | 1 |
| 1 | 1 | 1 |



Original

Filter

Result

(blur with a box filter)

Source: Bob Woodham, UBC. CPSC 425: Computer Vision lecture notes

Correlation and Convolution

Correlation (previous slides)

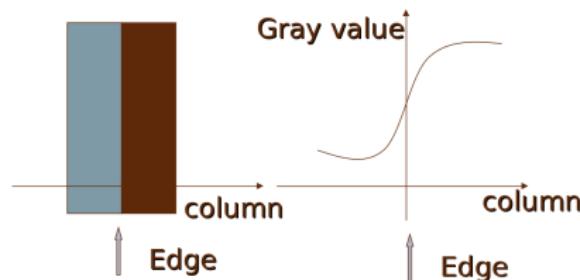
$$I'(X, Y) = \sum_{j=-k}^k \sum_{i=-k}^k F(i, j) I(X + i, Y + j)$$

Convolution

$$\begin{aligned} I'(X, Y) &= \sum_{j=-k}^k \sum_{i=-k}^k F(i, j) I(X - i, Y - j) \\ &= \sum_{j=-k}^k \sum_{i=-k}^k F(-i, -j) I(X + i, Y + j) \end{aligned}$$

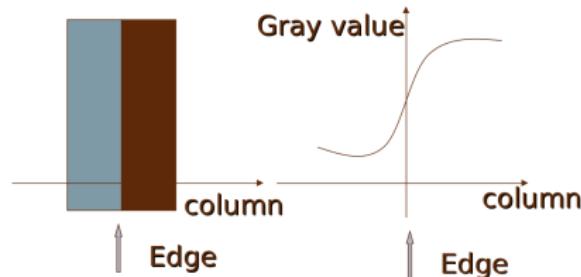
Note: If $F(X, Y) = F(-X, -Y)$ then correlation = convolution
i.e. if $F(X, Y)$ symmetric

Edges in Images

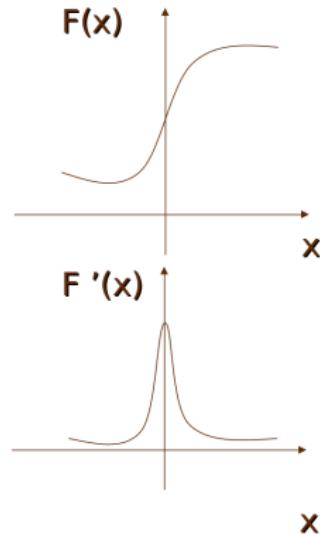


Edge = Sharp variation in images

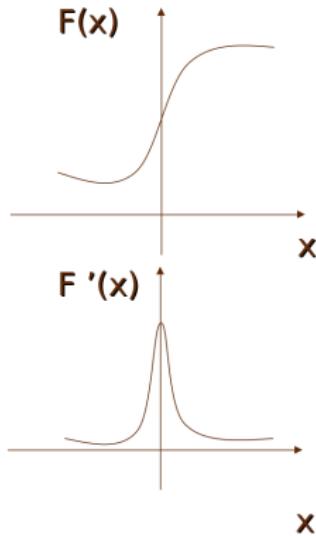
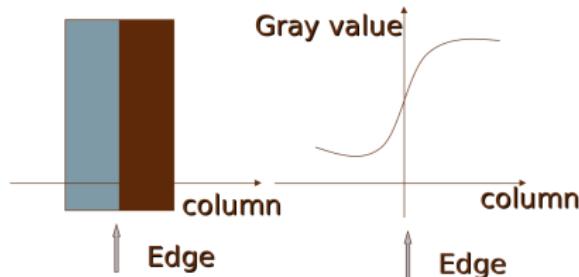
Edges in Images



Edge = Sharp variation in images
Edge \Rightarrow Large first derivative



Edges in Images



Edge = Sharp variation in images
Edge \Rightarrow Large first derivative

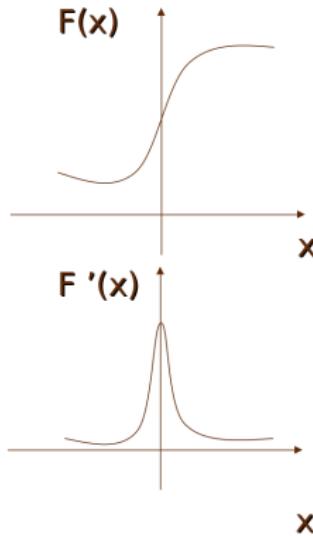
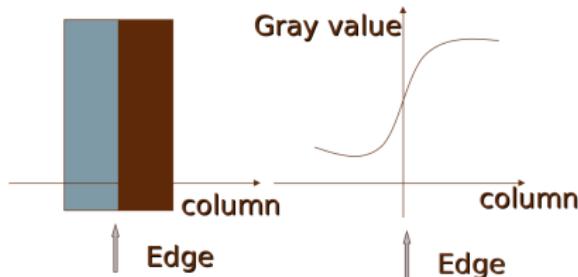
$$\frac{df(x)}{dx} = \lim_{\Delta x \rightarrow 0} \frac{f(x + \Delta x) - f(x)}{\Delta x}$$

$$\frac{df(x)}{dx} \cong \frac{f(x+1) - f(x-1)}{2}$$

Convolve with $\boxed{-1 \ 0 \ 1}$

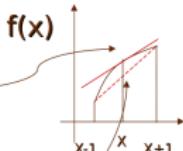
Source: Jana Kosecka, Stanford CS223b Computer Vision, Winter 2007

Edges in Images



Edge = Sharp variation in images
Edge \Rightarrow Large first derivative

$$\frac{df(x)}{dx} = \lim_{\Delta x \rightarrow 0} \frac{f(x + \Delta x) - f(x)}{\Delta x}$$



$$\frac{df(x)}{dx} \cong \frac{f(x+1) - f(x-1)}{2}$$

Convolve with -1 0 1

Vertical edges convolve with

| | | |
|----|---|---|
| -1 | 0 | 1 |
|----|---|---|

Horizontal edges convolve with

| |
|----|
| -1 |
| 0 |
| 1 |

Sobel Edge Detection



Input image I



Output image

Davidwkennedy,

<http://en.wikipedia.org/wiki/File:Bikesgray.jpg>

Sobel Edge Detection

Vertical edge detection

$$I \star \begin{array}{|c|c|c|} \hline -1 & 0 & 1 \\ \hline -2 & 0 & 2 \\ \hline -1 & 0 & 1 \\ \hline \end{array} \text{ gives } \frac{\partial I}{\partial x} = I_x$$

Horizontal edge detection:

$$I \star \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 0 & 0 & 0 \\ \hline -1 & -2 & -1 \\ \hline \end{array} \text{ gives } \frac{\partial I}{\partial y} = I_y$$



Input image I



Output image

Davidwkennedy,

<http://en.wikipedia.org/wiki/File:Bikesgray.jpg>

Sobel Edge Detection

Vertical edge detection

$$I * \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix} \text{ gives } \frac{\partial I}{\partial x} = I_x$$

Horizontal edge detection:

$$I * \begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix} \text{ gives } \frac{\partial I}{\partial y} = I_y$$

Each I_x and I_y is an image.

The image gradient:

$$\nabla I = \begin{bmatrix} I_x \\ I_y \end{bmatrix}$$

Magnitude of gradient image:

$$\|\nabla I\| = \sqrt{I_x^2 + I_y^2}$$



Input image I



Output image $\|\nabla I\|$

Davidwkennedy,

<http://en.wikipedia.org/wiki/File:Bikesgray.jpg>

Feature Extraction

Two types of feature

- ▶ Hand Engineered

Feature Extraction

Two types of feature

- ▶ Hand Engineered
 - ▶ easy to understand

Feature Extraction

Two types of feature

- ▶ Hand Engineered
 - ▶ easy to understand
 - ▶ if designed well, delivers good performance with small computation overhead

Feature Extraction

Two types of feature

- ▶ Hand Engineered
 - ▶ easy to understand
 - ▶ if designed well, delivers good performance with small computation overhead
 - ▶ requires a lot of expertise and domain knowledge

Feature Extraction

Two types of feature

- ▶ Hand Engineered
 - ▶ easy to understand
 - ▶ if designed well, delivers good performance with small computation overhead
 - ▶ requires a lot of expertise and domain knowledge
- ▶ Learned from data

Feature Extraction

Two types of feature

- ▶ Hand Engineered
 - ▶ easy to understand
 - ▶ if designed well, delivers good performance with small computation overhead
 - ▶ requires a lot of expertise and domain knowledge
- ▶ Learned from data
 - ▶ quite popular these days (Why ?)

Feature Extraction

Two types of feature

- ▶ Hand Engineered
 - ▶ easy to understand
 - ▶ if designed well, delivers good performance with small computation overhead
 - ▶ requires a lot of expertise and domain knowledge
- ▶ Learned from data
 - ▶ quite popular these days (Why ?)
 - ▶ requires a **lot of data** and **computational might**

Feature Extraction

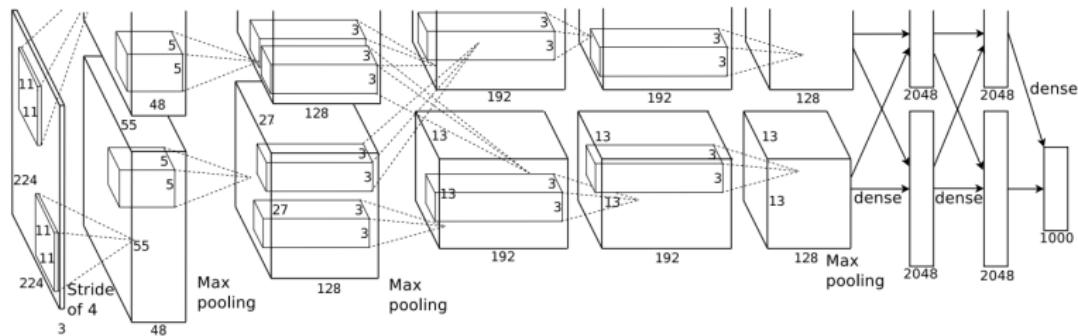
Two types of feature

- ▶ Hand Engineered
 - ▶ easy to understand
 - ▶ if designed well, delivers good performance with small computation overhead
 - ▶ requires a lot of expertise and domain knowledge
- ▶ Learned from data
 - ▶ quite popular these days (Why ?)
 - ▶ requires a **lot of data and computational might**
 - ▶ difficult (and sometime impossible) to understand (**a black box**)

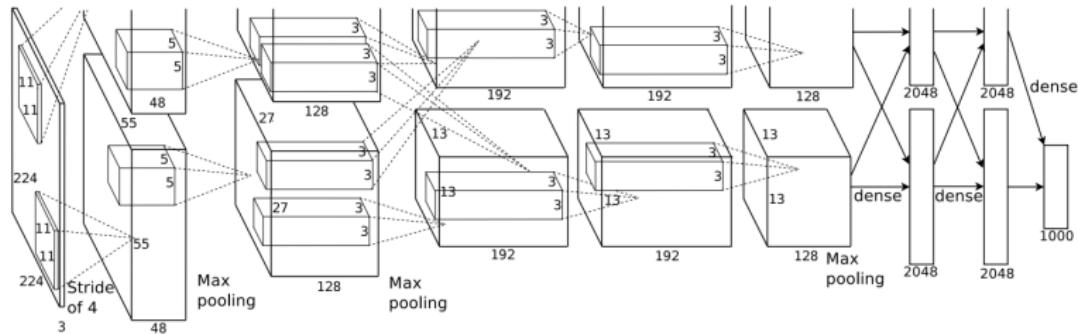
Feature Extraction : Learned from data

- ▶ Convolution Neural Network
- ▶ Auto-encoder
- ▶ ...

Feature Extraction : Learned from data - CNN

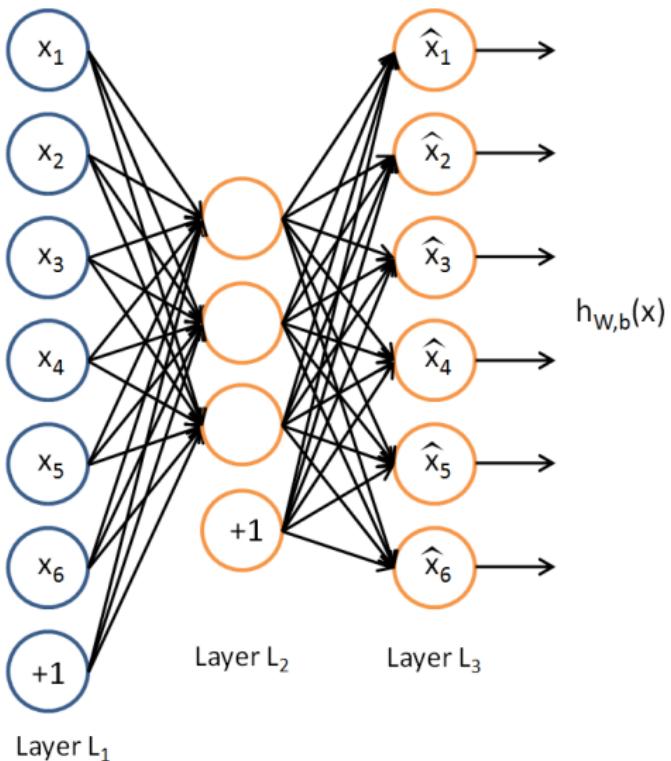


Feature Extraction : Learned from data - CNN



source: <http://colah.github.io/posts/2014-07-Conv-Nets-Modular/>

Feature Extraction : Learned from data - Autoencoder



Thank You